

137_Improving_Model_Utility_in.pdf

 Institute of Electrical and Electronics Engineers (IEEE)

Document Details

Submission ID

trn:oid:::14348:562137899

Submission Date

Mar 1, 2026, 10:01 AM GMT+8

Download Date

Mar 1, 2026, 11:21 AM GMT+8

File Name

137_Improving_Model_Utility_in.pdf

File Size

736.5 KB

8 Pages

6,399 Words

37,610 Characters

33% Overall Similarity

The combined total of all matches, including overlapping sources, for each database.

Match Groups

-  **103 Not Cited or Quoted 29%**
Matches with neither in-text citation nor quotation marks
-  **16 Missing Quotations 2%**
Matches that are still very similar to source material
-  **7 Missing Citation 2%**
Matches that have quotation marks, but no in-text citation
-  **0 Cited and Quoted 0%**
Matches with in-text citation present, but no quotation marks

Top Sources

- 28%  Internet sources
- 31%  Publications
- 1%  Submitted works (Student Papers)

Integrity Flags

1 Integrity Flag for Review

-  **Replaced Characters**
30 suspect characters on 5 pages
Letters are swapped with similar characters from another alphabet.

Our system's algorithms look deeply at a document for any inconsistencies that would set it apart from a normal submission. If we notice something strange, we flag it for you to review.

A Flag is not necessarily an indicator of a problem. However, we'd recommend you focus your attention there for further review.

Match Groups

-  **103 Not Cited or Quoted 29%**
Matches with neither in-text citation nor quotation marks
-  **16 Missing Quotations 2%**
Matches that are still very similar to source material
-  **7 Missing Citation 2%**
Matches that have quotation marks, but no in-text citation
-  **0 Cited and Quoted 0%**
Matches with in-text citation present, but no quotation marks

Top Sources

- 28%  Internet sources
- 31%  Publications
- 1%  Submitted works (Student Papers)

Top Sources

The sources with the highest number of matches within the submission. Overlapping sources will not be displayed.

1	Internet	arxiv.org	6%
2	Internet	vuir.vu.edu.au	6%
3	Internet	export.arxiv.org	3%
4	Internet	dora.dmu.ac.uk	1%
5	Internet	theses.gla.ac.uk	<1%
6	Internet	online-journals.org	<1%
7	Internet	research.usq.edu.au	<1%
8	Publication	Muhammad Ayat Hidayat, Yugo Nakamura, Yutaka Arakawa. "Privacy-Preserving..."	<1%
9	Internet	www.ijrte.org	<1%
10	Internet	www.medrxiv.org	<1%

11	Internet	1croreprojects.com	<1%
12	Publication	Moli Lyu, Zhiwei Ni, Qian Chen, Fenggang Li. "Edge-DPSDG: An Edge-based Differe...	<1%
13	Internet	d-nb.info	<1%
14	Publication	"Information Security and Privacy", Springer Science and Business Media LLC, 2025	<1%
15	Publication	Shekha Chenthara, Hua Wang, Khandakar Ahmed, Frank Whittaker, Ke Ji. "A Bloc...	<1%
16	Publication	Yuwei Sun, Hideya Ochiai, Hiroshi Esaki. "Decentralized Deep Learning for Multi-A...	<1%
17	Internet	fenix.tecnico.ulisboa.pt	<1%
18	Internet	poster-openaccess.com	<1%
19	Publication	Krishnan, Aditya. "Computationally Efficient, Privacy-Preserving, and Resource-Co...	<1%
20	Internet	eudl.eu	<1%
21	Internet	scholar.archive.org	<1%
22	Internet	proceedings.mlr.press	<1%
23	Publication	Zijuan Han, Yang Yang, Jinkai Zhang, Yunxia Liu, Ngai-Fong Bonnie Law. "DWT-RF...	<1%
24	Internet	www.mdpi.com	<1%

25	Internet	d197for5662m48.cloudfront.net	<1%
26	Publication	Enamul Kabir, Lili Sun, Hua Wang. "Finding an Optimum Set of Roles in a CPAC M...	<1%
27	Publication	Khadija Sultana, Khandakar Ahmed, Bruce Gu, Hua Wang. "Elastic Optimized Edg...	<1%
28	Internet	qmro.qmul.ac.uk	<1%
29	Publication	Lujin Zhao, Yi Lin, Sujuan Qin, Wenmin Li, Fei Gao, Yijie Shi, Zhengping Jin. "FedSC...	<1%
30	Publication	S.S. Iyengar, Seyedsina Nabavirazavi, Yashas Hariprasad, Prasad HB, C. Krishna M...	<1%
31	Internet	ia802506.us.archive.org	<1%
32	Publication	"Algorithms and Architectures for Parallel Processing", Springer Science and Busi...	<1%
33	Publication	Hakeem Quadri. "Mixed Bayesian Stackelberg Strategies for Robust Adversarial C...	<1%
34	Publication	Sidra Fareed, Ding Yi, Babar Hussain, Subhan Uddin, Aqsa Arif, Amir Nazar Tajoor...	<1%
35	Internet	theses.hal.science	<1%
36	Publication	"Advanced Intelligent Computing Technology and Applications", Springer Science...	<1%
37	Internet	link.springer.com	<1%
38	Publication	Assem Utaliyeva, Yoon-Ho Choi. "A user-centric privacy transformation framewor...	<1%

39	Publication	Basharat, Alina. "Towards Trustworthy Federated Learning.", The University of Te...	<1%
40	Publication	Chen Zhao, Jianping Hua, Michael L. Bittner, Ivan Ivanov, and Edward R. Dougher...	<1%
41	Publication	Li, Ao. "Tackling Data and Resource Heterogeneity for Performance Enhancement...	<1%
42	Publication	Runhua Xu, Nathalie Baracaldo, Yi Zhou, Ali Anwar, Heiko Ludwig. "HybridAlpha",...	<1%
43	Publication	Siyu Lu, Zheng Liu, Tianlin Liu, Wangchunshu Zhou. "Scaling-up medical vision-an...	<1%
44	Publication	Zhang, Zhenxiao. "Advancing Secure and Efficient Edge Intelligence Over Mobile ...	<1%
45	Publication	Ji Wang, Bokai Cao, Philip Yu, Lichao Sun, Weidong Bao, Xiaomin Zhu. "Deep Lear...	<1%
46	Publication	Kianidehkordi, Shahrzad. "Distributed Optimization Algorithms With Improved Ef...	<1%
47	Publication	Xiaoshi Wang, Haiyan Kang. "An enhanced method for hierarchical federated lear...	<1%
48	Internet	nowpublishers.com	<1%
49	Internet	repository.kaust.edu.sa	<1%
50	Publication	"Proceedings of International Conference on Computing Systems and Intelligent ...	<1%
51	Publication	Miao Du, Peng Yang, Yinqiu Liu, Xiaoming He, Mingkai Chen. "DP-Fed6G: An adap...	<1%
52	Publication	Muhammad Bilal Yaseen, Fayu Wan, Muhammad Bilal Riaz, Md Owahedur Rahma...	<1%

53	Publication	Thallita G.L. Lima, Wesley L. Passos. "When decentralization is not enough: Feder...	<1%
54	Publication	Wenqi Wei, Ling Liu, Jingya Zhou, Ka-Ho Chow, Yanzhao Wu. "Securing Distribute...	<1%
55	Internet	openaccess.city.ac.uk	<1%
56	Internet	repositorio.ufc.br	<1%
57	Internet	smartech.gatech.edu	<1%
58	Publication	Cho, Yae Jee. "Faster, Incentivized, and Efficient Federated Learning: Theory and ...	<1%
59	Publication	Fan, Xiaofeng. "Enhancing Sample Efficiency and Privacy in Real-World Decision-...	<1%
60	Publication	Martin Abadi, Andy Chu, Ian Goodfellow, H. Brendan McMahan, Ilya Mironov, Kun...	<1%
61	Publication	Ruiqi Zheng, Liang Qu, Guanhua Ye, Tong Chen, Yuhui Shi, Hongzhi Yin. "DeckKG: ...	<1%
62	Publication	Samsad Jahan, Yong-Feng Ge, Enamul Kabir, Kate Wang. "Analysis and Multi-obje...	<1%

Improving Model Utility in Privacy-Preserving Federated Learning

Meshari Alkhathami

Institute for Sustainable and

Liveable Cities,

Victoria University,

Melbourne 3011, Australia

meshari.alkhathami@live.vu.edu.au

Yong-Feng Ge

Institute for Sustainable and

Liveable Cities,

Victoria University,

Melbourne 3011, Australia

yongfeng.ge@vu.edu.au

Hua Wang

Institute for Sustainable and

Liveable Cities,

Victoria University,

Melbourne 3011, Australia

hua.wang@vu.edu.au

Abstract—Federated learning (FL) has emerged as a promising paradigm for collaborative model training while preserving data privacy by keeping raw data localized on client devices. To provide formal privacy guarantees, differential privacy (DP) is commonly integrated into FL frameworks; however, uniform noise injection across clients often results in significant degradation of model utility, especially under non-independent and identically distributed (non-IID) data conditions. This limitation highlights the need for privacy mechanisms that adapt to the diverse behaviors of clients in federated environments in the real-world. This paper introduces a novel adaptive client-level noise control mechanism that dynamically calibrates the noise added to each client's model update based on its divergence from the global model. By tailoring the noise scale to individual client behavior, the proposed method enhances the trade-off between privacy and utility while ensuring rigorous client-level DP. Empirical evaluations conducted on benchmark datasets demonstrate that the approach significantly outperforms existing baselines in terms of both accuracy and convergence speed. The results highlight the effectiveness of adaptive privacy mechanisms in addressing the challenges of utility loss in privacy-preserving FL.

Index Terms—Federated learning, Differential privacy, Adaptive privacy budget, Client-level privacy, Privacy-utility trade-off.

I. INTRODUCTION

Federated learning (FL) enables multiple clients to collaboratively train a machine learning model without directly sharing their private data [1]. This decentralized paradigm has emerged as a promising solution for privacy-sensitive applications such as mobile health, finance, and smart devices [2]–[4]. However, despite its inherent privacy advantages, FL remains susceptible to various privacy threats, particularly inference attacks such as model inversion, membership inference, and gradient leakage that exploit information embedded in local model updates [5]–[8].

Differential privacy (DP) has been widely adopted as a formal mechanism to mitigate these risks by injecting noise into the training process [9], [10]. Most existing DP-FL approaches enforce a uniform privacy budget across all clients [11], [12], applying the same level of noise regardless of individual client behavior or data characteristics. While such approaches provide strong worst-case privacy guarantees, they often degrade model performance, especially in heterogeneous (non-

IID) data environments, which are an inherent property of real-world FL deployments [13]–[15]. This tension between strong privacy guarantees and practical model utility has motivated a growing body of research on privacy-utility trade-offs in privacy-preserving learning systems [16]–[19].

This paper proposes a novel *adaptive client-level noise control* mechanism for FL. The method dynamically calibrates the noise injected into each client's update according to its divergence from the global training trend. Clients with updates that deviate significantly from the average, indicating potentially unique or sensitive data, receive stronger noise for enhanced privacy protection. In contrast, clients whose updates closely align with the population are assigned lower noise levels, thereby contributing higher-fidelity updates. This divergence-aware adaptation improves the overall privacy-utility trade-off in FL.

The main contributions of this paper are summarized as follows:

- A personalization-aware DP mechanism for FL that dynamically adapts the noise scale for each client based on its update divergence.
- A complete algorithmic framework that integrates divergence computation, noise scaling, ℓ_2 clipping, and aggregation within the FL pipeline.
- An empirical evaluation on benchmark datasets (MNIST and FEMNIST) under non-IID conditions, demonstrating the effectiveness of the proposed method in enhancing model utility while preserving privacy.

The remainder of this paper is organized as follows. Section II presents a comprehensive review of related work on federated learning and differential privacy, highlighting the key limitations and challenges that motivate this study. Section III introduces the proposed divergence-based adaptive client-level noise control mechanism and describes its integration within the federated learning framework. Section IV details the experimental design, datasets, and evaluation protocol used to assess the proposed approach. Section V reports and analyzes the empirical results, with particular emphasis on privacy-utility trade-offs. Section VI compares the proposed method with state-of-the-art differentially private federated

learning approaches. Finally, Section VII concludes the paper and outlines directions for future research.

II. RELATED WORK

Federated learning (FL) has emerged as a powerful paradigm for decentralized model training, enabling multiple clients to collaboratively learn a global model while keeping raw data localized [20]–[23]. This approach is particularly attractive for privacy-sensitive applications such as healthcare, mobile computing, and finance [24]–[27]. However, despite its decentralized architecture, FL is not immune to privacy risks. Prior research has demonstrated that gradients or model updates can leak sensitive information through reconstruction attacks, gradient inversion, and membership inference [28], [29]. These vulnerabilities underscore the need to augment FL with formal privacy-preserving mechanisms, among which differential privacy has emerged as the most prominent and widely adopted framework [9].

Client-Level Differential Privacy

The introduction of DP-SGD by Abadi et al. [30] demonstrated how DP can be combined with deep learning in centralized settings by clipping gradients and injecting Gaussian noise. Building upon this foundation, Geyer et al. [11] proposed one of the earliest frameworks for client-level DP in FL. Their method clips each client's update and applies Gaussian noise before aggregation, thus protecting entire client datasets. They further employed the moments accountant [30] to track cumulative privacy loss. While influential, this line of work adopts a uniform noise scale across all clients, disregarding data heterogeneity. As a result, clients with highly representative updates may suffer unnecessary utility loss, whereas clients with atypical data remain at higher risk under insufficient noise. McMahan et al. [12] extended this idea through the DP-FedAvg algorithm, which integrates secure aggregation [31] with DP accounting. Nevertheless, these methods remain limited by their reliance on fixed, global privacy budgets that treat all clients equally regardless of contribution diversity.

Personalized Differential Privacy Mechanisms

Recognizing the limitations of uniform noise, subsequent research has introduced personalization into DP-FL. Wei et al. [32] proposed privacy accounting strategies that adapt noise levels according to participation frequency, offering stronger utility guarantees for active clients. Triastcyn and Faltings [33] introduced a Bayesian framework for personalized DP, allowing clients to calibrate privacy guarantees based on probabilistic sensitivity modeling. Li et al. [34] advanced clustering-based approaches, grouping clients with similar data distributions and assigning differentiated noise budgets. More broadly, adaptive evolutionary frameworks have been proposed to support dynamic privacy-preserving data publishing under changing data distributions, highlighting the importance of adaptivity beyond static personalization [35]. These efforts demonstrated that personalization can mitigate the over-penalization of uniform DP. However, most approaches rely

on static heuristics (such as participation rates, metadata, or pre-defined clusters) [34], which may not adequately reflect real-time client behavior or update sensitivity. This lack of adaptivity reduces their robustness in dynamic, large-scale FL environments [15], [36].

Adaptive and Utility-Aware Privacy Strategies

Beyond static personalization, researchers have begun to explore adaptive mechanisms that dynamically tune privacy parameters during training. Yang et al. [15] proposed dynamic personalized FL with adaptive DP, where privacy budgets are adjusted during training to better balance privacy and utility. Sun et al. [37] introduced LDP-FL, a framework for practical private aggregation in FL using local differential privacy. Andrew et al. [36] developed adaptive clipping, which minimizes sensitivity prior to noise injection and improves convergence under DP constraints. Fu et al. [38] studied adaptive DP mechanisms in FL through per-client, per-round gradient clipping and varying noise scales, showing that such adaptation can improve convergence and accuracy under heterogeneous data. Related optimization-based approaches have also been explored in privacy-preserving data publishing to explicitly balance privacy and utility objectives [39]. Despite these advances, many adaptive strategies still rely primarily on local gradient norms or pairwise similarity, rather than aligning updates semantically with the global model.

As a result, such methods may overlook cases where a client's divergence from the global trend signals both contribution uniqueness and elevated privacy risk. Mireshghallah et al. [40] and Kairouz et al. [13] further emphasized the difficulty of reconciling strong theoretical DP guarantees with practical utility, noting that many adaptive mechanisms lack comprehensive validation under realistic, large-scale non-IID scenarios. In addition to gradient-based approaches, federated evolutionary optimization techniques have also been explored for privacy-preserving learning and data publishing in distributed environments [41].

Limitations of Existing Work

Across these three research directions, several common limitations remain. First, early client-level DP methods [11], [20] enforce uniform noise across all clients, which fails to account for heterogeneity and results in significant utility degradation. Second, personalized methods [32]–[34] introduce flexibility but rely on static heuristics or clustering, which may not adapt to changing training dynamics. Third, adaptive strategies [36]–[38], [42] incorporate gradient norms or similarity measures but often ignore the divergence of client updates from the global model trajectory—a critical factor for capturing atypical or sensitive contributions. Finally, much of the existing literature emphasizes formal guarantees while offering limited empirical evaluation under realistic, large-scale non-IID conditions, leaving open questions about their practical utility and scalability [13].

This paper proposes a divergence-based adaptive client-level DP mechanism that dynamically calibrates noise in each

client's model update. It uses the deviation of local updates from the global mean to assess contribution atypicality and privacy risk. Unlike existing methods that apply uniform noise or static heuristics [43]–[45], our framework aligns privacy protection with each client's alignment to the global learning trajectory. Updates with higher divergence, often linked to less representative or sensitive data, receive stronger perturbation, while well-aligned updates get milder noise. This approach reduces unnecessary perturbation for most clients while ensuring robust privacy for atypical contributors. Experiments on MNIST and FEMNIST under non-IID conditions show that our method achieves a better privacy–utility trade-off, improving accuracy and stability compared to fixed-noise and non-adaptive baselines.

III. PROPOSED ALGORITHM

This section outlines the proposed approach for achieving personalization-aware client-level DP in FL. Unlike prior methods that enforce a uniform privacy budget and apply equal noise to all clients, the proposed method introduces an adaptive mechanism that dynamically adjusts the noise injected into each client's update based on its deviation from the global update trend. This approach is designed to enhance privacy protection for clients with atypical data distributions while preserving model utility for clients whose data aligns more closely with the overall population.

A. Preliminaries and Notation

FL is a distributed training paradigm in which a central server coordinates the collaborative optimization of a model across a population of clients. At communication round t , the server maintains the global model w^t . Each selected client k downloads w^t , performs local training on its private dataset D_k , and obtains a local model w_k . The contribution of client k is expressed as a model update:

$$\Delta w_k = w_k - w^t,$$

which reflects the modification induced by local training relative to the global model.

Before aggregation, two standard mechanisms are applied to ensure both fairness and privacy:

- **Update clipping:** limits the maximum influence of any individual client by constraining the ℓ_2 -norm of updates.
- **Noise injection:** perturbs clipped updates with Gaussian noise, thereby providing formal DP guarantees against potential inference attacks.

These definitions and operations establish the foundation for the proposed adaptive noise control method.

B. Federated Learning Setup

Consider a cross-device FL setting involving K clients. In each communication round t , a subset $Z_t \subseteq \{1, \dots, K\}$ of m_t clients is sampled uniformly at random to participate. Each selected client $k \in Z_t$ receives the current global model w_t and

performs local training on its private dataset D_k . The resulting local update is computed as:

$$\Delta w_k = w_k - w_t,$$

where w_k denotes the locally trained model at client k .

C. Update Clipping

To bound the sensitivity of individual updates and control the influence of any single client, each local update is clipped using an ℓ_2 -norm threshold C . This is expressed as:

$$\tilde{\Delta w}_k = \frac{\Delta w_k}{\max\left(1, \frac{\|\Delta w_k\|_2}{C}\right)}.$$

This operation ensures that all updates lie within a Euclidean ball of radius C , thereby standardizing the scale of client contributions prior to the addition of noise.

D. Divergence-Based Noise Scaling

In contrast to existing methods that apply a uniform noise level to all clients [11], we introduce a divergence-aware mechanism that adapts the magnitude of noise injected into each client's model update based on its deviation from the round-wise average behavior. This approach aims to provide stronger privacy protection for atypical clients—those whose updates significantly deviate from the norm—while preserving the utility of clients that align more closely with the global learning trend.

Let $\tilde{\Delta w}_k$ denote the ℓ_2 -clipped model update from client k in communication round t . We compute the mean update across all participating clients as:

$$\mu_t = \frac{1}{m_t} \sum_{k \in Z_t} \tilde{\Delta w}_k,$$

where Z_t is the set of m_t clients sampled in round t .

We then define the divergence of each client's update as the Euclidean distance from the mean:

$$D_k = \left\| \tilde{\Delta w}_k - \mu_t \right\|_2.$$

This divergence score D_k quantifies how much a client's contribution differs from the average client behavior. A larger divergence indicates a potentially more unique or sensitive local dataset, warranting stronger perturbation.

To personalize the noise level accordingly, we define the client-specific noise scale as:

$$\sigma_k = \sigma_{\text{base}} + \alpha \cdot D_k,$$

where σ_{base} is a global base noise multiplier and α is a tunable sensitivity coefficient. This linear scaling ensures that clients with greater divergence receive proportionally more noise, while low-divergence clients are less affected, thereby reducing unnecessary utility loss.

We note that σ_k can optionally be bounded by a maximum noise cap to prevent instability. In practice, we find this divergence-aware scaling enables better alignment between privacy needs and contribution representativeness, especially in non-IID data settings.

E. Noisy Update Aggregation

After computing the client-specific noise scales based on update divergence, each participating client injects Gaussian noise into its clipped model update to achieve DP. The noise added to each client's update is scaled according to its personalized noise level σ_k , determined as described in the previous section.

Each client k produces a differentially private update as:

$$\hat{\Delta}w_k = \tilde{\Delta}w_k + \mathcal{N}(0, \sigma_k^2 C^2 I),$$

where $\mathcal{N}(0, \sigma_k^2 S^2 I)$ is multivariate Gaussian noise with zero mean and covariance proportional to both the personalized noise scale and the clipping threshold C .

This formulation ensures that all updates lie within a bounded sensitivity range (due to clipping) and are appropriately perturbed to limit the influence of any individual client. The noise is applied independently to each client, allowing for fine-grained privacy control across the network.

Once all noisy updates $\hat{\Delta}w_k$ are received, the server performs standard federated aggregation by averaging the contributions from all m_t participating clients:

$$w_{t+1} = w_t + \frac{1}{m_t} \sum_{k \in \mathcal{Z}_t} \hat{\Delta}w_k.$$

This aggregation procedure maintains unbiased updates while incorporating per-client DP through personalized perturbation. By adapting the amount of noise each client adds based on its divergence, the method effectively reduces the privacy-utility trade-off that arises in uniform-noise approaches, particularly in non-IID federated settings.

F. Adaptive Client-Level DP Mechanism

Client-level (ϵ, δ) -DP is enforced by adding Gaussian noise after clipping, with a client-specific scale σ_k based on divergence from the round mean:

$$\bar{u} = \frac{1}{K} \sum_{k=1}^K \tilde{\Delta}w_k, \quad (1)$$

$$d_k = \left\| \tilde{\Delta}w_k - \bar{u} \right\|_2, \quad (2)$$

$$\hat{d}_k = \min(d_k, 1). \quad (3)$$

The adaptive noise scale is defined as

$$\sigma_k = \begin{cases} 0, & t \leq T_{\text{warmup}}, \\ \min(\sigma_{\max}, \sigma_{\text{base}} + \alpha \hat{d}_k), & t > T_{\text{warmup}}, \end{cases} \quad (4)$$

$$\hat{\Delta}w_k = \tilde{\Delta}w_k + \mathcal{N}(0, \sigma_k^2 C^2 I). \quad (5)$$

Intuition: atypical (high-divergence) updates receive stronger perturbation (up to $\sigma_{\max}$), while well-aligned updates receive milder noise, reducing unnecessary utility loss. A warm-up phase ($t \leq T_{\text{warmup}}$) disables noise injection to stabilize early representation learning under severe heterogeneity; as no noise is applied, these rounds are excluded from privacy accounting and client-level DP guarantees apply only after noise activation.

G. Privacy Accounting

Due to the heterogeneity in client-specific noise levels, the resulting privacy guarantees vary across participants. To track the cumulative privacy loss over successive communication rounds, we adopt the moments accountant framework [30], which provides a tight and composable analysis of privacy expenditure under repeated randomized mechanisms.

In the proposed framework, each client k is associated with a potentially different noise scale σ_k , determined adaptively based on update divergence. As a result, the privacy loss incurred by each client depends on both its noise level and its frequency of participation. The moments accountant enables accurate tracking of individual privacy parameters (ϵ_k, δ_k) by accumulating the privacy loss incurred at each communication round, while accounting for subsampling and Gaussian perturbation.

This per-client accounting is particularly important in adaptive settings, where privacy expenditure may evolve differently across clients. By monitoring individual privacy budgets throughout training, the system can enforce predefined privacy constraints. In practice, client participation may be halted or limited once a target privacy budget is exhausted, ensuring compliance with client-level privacy guarantees.

H. Overall Procedure

This section presents the complete procedure of the proposed personalized FL algorithm with adaptive client-level DP. The method dynamically calibrates the noise added to each client's local model update based on its deviation from the average update. This design provides stronger privacy protection for clients with atypical or sensitive data distributions while maintaining high model utility for the majority of participants.

Algorithm 1 Adaptive Client-Level DP in FL

Require: Global model w_0 , total rounds T , number of clients K , base noise σ_{base} , scaling factor α , clipping threshold S , learning rate η

Ensure: Global model after final communication round w_T

```

0: for  $t = 1$  to  $T$  do
0:   Sample a subset of clients  $\mathcal{C}_t \subseteq \{1, \dots, K\}$ 
0:   for all client  $k \in \mathcal{C}_t$  in parallel do
0:     Receive global model  $w_t$  from the server
0:     Compute local update:  $\Delta w_k = \text{ClientUpdate}(w_t, \mathcal{D}_k)$ 
0:     Clip update:  $\tilde{\Delta}w_k = \Delta w_k / \max(1, \|\Delta w_k\|_2 / S)$ 
0:   end for
0:   Compute average update:  $\bar{\Delta}w = \frac{1}{|\mathcal{C}_t|} \sum_{k \in \mathcal{C}_t} \tilde{\Delta}w_k$ 
0:   for all client  $k \in \mathcal{C}_t$  do
0:     Compute divergence:  $d_k = \|\tilde{\Delta}w_k - \bar{\Delta}w\|_2$ 
0:     Compute noise scale:  $\sigma_k = \min(\sigma_{\max}, \sigma_{\text{base}} + \alpha \cdot d_k)$ 
0:     Add noise:  $\hat{\Delta}w_k \leftarrow \tilde{\Delta}w_k + \mathcal{N}(0, \sigma_k^2 S^2 I)$ 
0:   end for
0:   Aggregate global model:  $w_{t+1} = w_t + \frac{1}{|\mathcal{C}_t|} \sum_{k \in \mathcal{C}_t} \hat{\Delta}w_k$ 
0:   Update privacy accountant
0: end for
1: return  $w_T = 0$ 

```

Algorithm 1 outlines the full training procedure. In each communication round t , a random subset of clients $\mathcal{C}_t$ is selected for participation. Each client k computes a local

43 model update Δw_k by training on its private dataset $\mathcal{D}_k$. To bound sensitivity, this update is clipped using an ℓ_2 -norm threshold S , yielding the clipped update $\tilde{\Delta w}_k$.

55 The server then computes the mean update $\bar{\Delta w}$ across all participating clients. To enable personalized privacy adaptation, the divergence $d_k = \|\tilde{\Delta w}_k - \bar{\Delta w}\|_2$ is calculated for each client. Based on this divergence, the personalized noise scale σ_k is computed as:

$$\sigma_k = \min(\sigma_{\max}, \sigma_{\text{base}} + \alpha \cdot d_k),$$

5 where σ_{base} is the base noise level, and α controls the adaptation strength. Gaussian noise sampled from $\mathcal{N}(0, \sigma_k^2 S^2 \mathbf{I})$ is then added to the clipped update to produce the differentially private update $\hat{\Delta w}_k$.

19 Finally, the server aggregates all private updates and updates 31 the global model. A privacy accountant tracks the cumulative privacy loss over rounds to ensure that each client remains within their allocated privacy budget. This adaptive framework achieves tighter privacy guarantees for sensitive clients while maintaining overall model performance.

1 IV. EXPERIMENTS

We evaluate the proposed *adaptive client-level DP* mechanism in a simulated cross-device FL setting. The design follows standard FL evaluation practice, reporting utility (accuracy, loss), training stability, and privacy-utility trends under non-IID data. We detail datasets, client partitioning, models, training protocol, and the adaptive noise mechanism for exact reproducibility.

37 A. Datasets and Client Partitioning

a) *MNIST*.: The dataset contains 60,000 training and 10,000 test grayscale digit images (28×28 , 10 classes). To simulate $N = 300$ clients, we partition the training set into 100 equal, non-overlapping shards (one shard per client), which induces mild label imbalance and non-IID behavior. Images are normalized to $[0, 1]$.

57 b) *FEMNIST*.: We use the public federated split (writers as clients) with 62 classes (digits and letters). To ensure stable local optimization, we select 100 writers each with ≥ 100 examples. Images are normalized to $[0, 1]$. A held-out test set is formed via a 90/10 split of the global training pool while preserving label coverage.

31 B. Model Architectures and Optimizers

a) *MNIST model*.: A compact MLP: Flatten(28×28) $\rightarrow$ Dense(128, ReLU) $\rightarrow$ Dense(10, softmax). We use Adam with sparse cross-entropy.

b) *FEMNIST model*.: A deeper CNN suited to the 62-class task and writer variability:

- Reshape ($28 \times 28 \times 1$)
- Conv2D(32, 3×3 , ReLU, same) + MaxPool(2×2)
- Conv2D(64, 3×3 , ReLU, same) + MaxPool(2×2)
- Conv2D(128, 3×3 , ReLU, same) + MaxPool(2×2)
- Flatten $\rightarrow$ Dense(512, ReLU) $\rightarrow$ Dropout(0.4) $\rightarrow$ Dense(62, softmax)

Optimization uses SGD with momentum 0.9 and exponential decay (initial LR 0.005, decay rate 0.95 every 1000 steps).

C. Federated Protocol

We adopt the standard Federated Averaging (FedAvg) framework with uniform client sampling. Let w_t denote the global model maintained by the central server at communication round t . In each round, a subset of K clients is sampled uniformly at random from the population of available clients. Each selected client receives the current global model w_t and performs local training for a fixed number of epochs on its private dataset. After local optimization, client k computes a model update

$$\Delta w_k = w_k - w_t,$$

where w_k denotes the locally updated model.

To bound the influence of any single client and control the sensitivity of the aggregation process, all client updates are clipped using an ℓ_2 -norm threshold C . Specifically, each update is normalized as

$$\tilde{\Delta w}_k = \frac{\Delta w_k}{\max\left(1, \frac{\|\Delta w_k\|_2}{C}\right)}.$$

This clipping operation ensures that the contribution of each client lies within a bounded region, which is a necessary condition for enforcing client-level privacy guarantees and maintaining numerical stability during aggregation.

Once the clipped updates are obtained, the server aggregates the received contributions by computing their average. The global model is then updated according to

$$w_{t+1} = w_t + \frac{1}{K} \sum_{k=1}^K \tilde{\Delta w}_k,$$

where $\tilde{\Delta w}_k$ denotes the (potentially perturbed) update returned by client k . In the proposed framework, noise is injected into $\tilde{\Delta w}_k$ prior to aggregation as part of the adaptive client-level privacy mechanism described in subsequent sections.

This protocol preserves the decentralized nature of FL while enabling fine-grained control over individual client contributions. Uniform sampling ensures fairness across clients, while the combination of update clipping and aggregation provides a robust foundation for integrating adaptive noise mechanisms. Importantly, this design remains compatible with standard FedAvg implementations and can be seamlessly integrated into existing FL pipelines without modifying the underlying optimization procedure.

D. Hyperparameters (Exact)

TABLE I
FEDERATED AND DP SETTINGS USED IN ALL EXPERIMENTS

Parameter	MNIST	FEMNIST
Total clients N	100 (simulated)	3,550 writers (≥ 100 samples)
Clients per round K	20	20
Rounds	300	800
Local epochs	5	15
Batch size	32	64
Clipping norm C	1.0	1.0
σ_{base}	0.01	0.01
σ_{max}	0.50	0.30
Adaptivity α	0.10	0.10
Warm-up T_{warmup}	0	50
Optimizer	Adam	SGD + momentum 0.9
LR schedule	constant (Adam)	exp. decay (init 0.005)

E. Evaluation Protocol

At the end of each communication round, the global model is evaluated on a held-out test set (standard MNIST test; FEMNIST test via a 90/10 split). We report top-1 accuracy and cross-entropy loss. To analyze privacy–utility dynamics, we also track the round-wise mean noise scale $\bar{\sigma}_k$. The training process is instrumented to record these evaluation quantities at every round to support reproducibility, facilitate quantitative comparisons, and enable subsequent statistical analysis.

For fair comparison, we reimplemented two baselines within the same training framework, as illustrated in Fig. 1: (i) **Static DP**, where client updates are clipped and perturbed with a fixed Gaussian noise scale across all clients [11]; and (ii) **No DP**, which corresponds to standard FedAvg without clipping or noise injection, serving as a non-private upper bound [20].

V. RESULTS AND DISCUSSION

A. Results

a) **MNIST**.: The adaptive mechanism retains strong utility under client-level DP, converging to approximately 93.5–94.0% test accuracy with low average noise ($\sigma_k \approx 0.107$). Convergence is faster and final utility higher than with fixed uniform noise at matched upper bounds, indicating that per-client calibration mitigates the typical uniform-noise penalty.

b) **FEMNIST**.: Accuracy rises quickly during the warm-up (noise-free) phase, peaking near $\sim 68\%$ top-1 accuracy; after enabling adaptive noise, performance stabilizes in the 55–60% range with $\sigma_k \approx 0.10$ –0.11. Training remains numerically stable (no NaNs) due to gradient clipping and bounded σ_k , while warm-up rounds do not contribute to privacy loss as noise injection is disabled.

c) **Summary**.: Across both datasets, adaptivity improves the privacy–utility balance relative to fixed-noise client-level DP: high-divergence clients receive proportionally more protection, while majority clients contribute higher-fidelity updates, yielding faster convergence and stronger final accuracy in practice. Importantly, the observed behavior in Fig. 2 illustrates that the adaptive mechanism preserves training stability across heterogeneous data regimes. On MNIST, the majority




Fig. 1. Test accuracy on MNIST across federated training rounds.

TABLE II
SUMMARY OF UTILITY UNDER ADAPTIVE CLIENT-LEVEL DP

Dataset	Rounds	Peak Acc. (%)	Final Acc. (%)	$\bar{\sigma}_k$ (post-warmup)
MNIST	300	93.5–94.0	93.5–94.0	≈ 0.107
FEMNIST	800	67.5–68.5 [†]	≈ 60	≈ 0.10 –0.11

[†]Peak occurred during warm-up (noise-free); after adaptive noise was enabled, accuracy stabilized with modest utility loss.

of clients rapidly align with the global model, resulting in consistently low average noise and near non-private accuracy.




Fig. 2. Comparison of adaptive DP training on MNIST (left) and FEMNIST (right). Top: test accuracy across rounds. Bottom: average client-level noise $\bar{\sigma}_k$ across rounds. Curves correspond to representative runs; variance analysis across multiple random seeds is deferred to future work.

B. Discussion

The empirical results (Fig. 2, left) demonstrate that adaptive client-level DP significantly improves the privacy–utility balance compared to uniform-noise baselines. On MNIST, the

method consistently maintained accuracy above 91%, with peak values around 92.5%. The average injected noise remained modest ($\bar{\sigma}_k \approx 0.107$), indicating that most clients contributed high-fidelity updates without excessive perturbation. These findings highlight that divergence-aware calibration can effectively preserve utility while still enforcing rigorous client-level DP.

The FEMNIST benchmark (Fig. 2, right), characterized by a larger label space and stronger heterogeneity across writers, presented a more demanding evaluation. During the warm-up stage, the model rapidly approached $\sim 68\%$ accuracy. Once adaptive noise was introduced, performance stabilized between 57% and 60%, avoiding collapse or instability even under severe non-IID conditions. The observed gap between peak and final accuracy illustrates the inherent privacy–utility trade-off, but the adaptive mechanism reduced unnecessary degradation relative to fixed-noise strategies by applying stronger perturbations only to atypical clients.

Overall, the results confirm that divergence-based noise assignment enhances stability and accelerates convergence, while hyperparameters α , σ_{base} , and C govern the privacy–utility trade-off by increasing protection for highly divergent clients at the cost of accuracy when overly aggressive.

VI. COMPARISON WITH STATE-OF-THE-ART METHODS

To further contextualize the effectiveness of the proposed method, we compare it with representative state-of-the-art DP-FL approaches.

As shown in Table III and Fig. 3, our approach achieves state-of-the-art performance on MNIST with a final accuracy of 93.5–94.0%, surpassing the previously reported results of 92–93% achieved by the dynamic personalized DP-FL approach. This demonstrates that our method not only provides stronger utility on MNIST but also delivers competitive performance on FEMNIST, where stable final accuracy of approximately 60% is achieved under adaptive noise.

The comparison highlights several important observations. First, classical client-level DP methods such as DP-FedAvg suffer from substantial utility degradation due to the use of a uniform noise scale across all clients, which fails to account for heterogeneity in data distributions. Second, adaptive clipping improves numerical stability but does not directly address the over-perturbation of representative client updates.

More recent adaptive and personalized DP-FL approaches achieve stronger utility by introducing client-specific mechanisms. However, many of these methods rely primarily on local gradient norms or explicit privacy budget allocation, which can increase algorithmic complexity and may not fully capture how individual updates align with the global learning objective.

VII. CONCLUSION

This paper proposed an adaptive client-level DP framework for federated learning that calibrates the noise added to each client update according to its divergence from the global model trajectory. By aligning privacy protection with the

TABLE III
COMPARISON WITH STATE-OF-THE-ART DP-FL METHODS

Method	Adaptivity	MNIST (%)	FEMNIST (%)
DP-FedAvg [11]	None	88–90	55–58
Adaptive Clipping [36]	Clip-level	90–91	–
AdaDP-FL [38]	Per-client	89–91	56–60
Dynamic Pers. DP-FL [15]	Per-client / round	92–93	60–63
Ours (Adaptive Noise)	Divergence-aware	93.5–94.0	68.0




Fig. 3. Comparison of final test accuracy under client-level DP for representative state-of-the-art FL methods on MNIST and FEMNIST.

atypicality of client contributions, the proposed approach reduces unnecessary perturbation for representative clients while providing stronger protection for those with more sensitive or unique local data. Experimental results demonstrate a favorable privacy–utility trade-off under non-IID data. On MNIST, the framework converges to approximately 93.5–94.0% test accuracy with low average noise, while on FEMNIST it maintains stable performance around 60% once adaptive noise is applied, despite severe heterogeneity. These results confirm that divergence-aware noise calibration effectively mitigates the utility degradation commonly observed in client-level DP-FL.

REFERENCES

- [1] Q. Yang, Y. Liu, T. Chen, and Y. Tong, “Federated machine learning: Concept and applications,” *ACM Transactions on Intelligent Systems and Technology (TIST)*, vol. 10, no. 2, pp. 1–19, 2019.
- [2] W. Y. B. Lim, N. C. Luong, D. T. Hoang, Y. Jiao, Y.-C. Liang, Q. Yang, D. Niyato, and C. Miao, “Federated learning in mobile edge networks: A comprehensive survey,” *IEEE communications surveys & tutorials*, vol. 22, no. 3, pp. 2031–2063, 2020.
- [3] D. Ning, Y.-F. Ge, H. Wang, and C. Zhou, “Cadif-osn: Detecting cloned accounts with missing profile attributes on online social networks,” in *Proceedings of the 33rd ACM International Conference on Information and Knowledge Management*, ser. CIKM ’24. ACM, Oct. 2024, pp. 1795–1803.
- [4] Y.-F. Ge, M. Orlowska, J. Cao, H. Wang, and Y. Zhang, “Mdde: multi-tasking distributed differential evolution for privacy-preserving database fragmentation,” *The VLDB Journal*, vol. 31, no. 5, pp. 957–975, Jan. 2022.
- [5] B. Hitaj, G. Ateniese, and F. Perez-Cruz, “Deep models under the gan: Information leakage from collaborative deep learning,” in *Proceedings of the 2017 ACM SIGSAC conference on computer and communications security*, 2017, pp. 603–618.
- [6] Y.-F. Ge, H. Wang, J. Cao, Y. Zhang, and X. Jiang, “Privacy-preserving data publishing: an information-driven distributed genetic algorithm,” *World Wide Web*, vol. 27, no. 1, Jan. 2024.

- [7] X. Sun, H. Wang, J. Li, and Y. Zhang, "Satisfying privacy requirements before data anonymization," *Comput. J.*, vol. 55, pp. 422–437, 2012. [Online]. Available: <https://api.semanticscholar.org/CorpusID:1315690>
- [8] X. Sun, M. Li, and H. Wang, "A family of enhanced (l, α)-diversity models for privacy preserving data publishing," *Future Gener. Comput. Syst.*, vol. 27, pp. 348–356, 2011. [Online]. Available: <https://api.semanticscholar.org/CorpusID:20012270>
- [9] C. Dwork, A. Roth *et al.*, "The algorithmic foundations of differential privacy," *Foundations and trends® in theoretical computer science*, vol. 9, no. 3–4, pp. 211–407, 2014.
- [10] S. Jahan, Y.-F. Ge, H. Wang, and E. Kabir, "Adaptive-parameter memetic algorithm for privacy-preserving trajectory data publishing: A multi-objective optimization approach," *Computing*, vol. 107, no. 7, Jun. 2025.
- [11] R. C. Geyer, T. Klein, and M. Nabi, "Differentially private federated learning: A client level perspective," *arXiv preprint arXiv:1712.07557*, 2017.
- [12] H. B. McMahan, D. Ramage, K. Talwar, and L. Zhang, "Learning differentially private recurrent language models," in *International Conference on Learning Representations (ICLR)*, 2018.
- [13] P. Kairouz, H. B. McMahan, B. Avent, A. Bellet, M. Bennis, A. N. Bhagoji, K. Bonawitz, Z. Charles, G. Cormode, R. Cummings *et al.*, "Advances and open problems in federated learning," *Foundations and trends® in machine learning*, vol. 14, no. 1–2, pp. 1–210, 2021.
- [14] Y.-F. Ge, H. Wang, E. Bertino, J. Cao, Y. Zhang, and Z. Zheng, "Distributed bandit-based cooperative coevolution for large-scale multi-objective data publishing," *IEEE Transactions on Services Computing*, 2025.
- [15] X. Yang, W. Huang, and M. Ye, "Dynamic personalized federated learning with adaptive differential privacy," *Advances in Neural Information Processing Systems*, vol. 36, pp. 72 181–72 192, 2023.
- [16] Y.-F. Ge, H. Wang, E. Bertino, J. Cao, and Y. Zhang, "Multiobjective privacy-preserving task assignment in spatial crowdsourcing," *IEEE Transactions on Cybernetics*, 2025.
- [17] Y.-F. Ge, H. Wang, E. Bertino, Z.-H. Zhan, J. Cao, Y. Zhang, and J. Zhang, "Evolutionary dynamic database partitioning optimization for privacy and utility," *IEEE Transactions on Dependable and Secure Computing*, vol. 21, no. 4, pp. 2296–2311, Jul. 2024.
- [18] Y.-F. Ge, E. Bertino, H. Wang, J. Cao, and Y. Zhang, "Distributed cooperative coevolution of data publishing privacy and transparency," *ACM Transactions on Knowledge Discovery from Data*, vol. 18, no. 1, pp. 1–23, Sep. 2023.
- [19] S. Jahan, Y.-F. Ge, E. Kabir, and K. Wang, "Analysis and multi-objective protection of public medical datasets from privacy and utility perspectives," *Data Science and Engineering*, vol. 10, no. 3, pp. 362–375, Apr. 2025.
- [20] B. McMahan, E. Moore, D. Ramage, S. Hampson, and B. A. y Arcas, "Communication-efficient learning of deep networks from decentralized data," in *Artificial intelligence and statistics*. PMLR, 2017, pp. 1273–1282.
- [21] M. You, Y.-F. Ge, K. Wang, H. Wang, J. Cao, and G. Kambourakis, "Hierarchical adaptive evolution framework for privacy-preserving data publishing," *World Wide Web*, vol. 27, no. 4, Jul. 2024.
- [22] M. E. Kabir and H. Wang, "Conditional purpose based access control model for privacy protection," in *Australasian Database Conference*, 2009. [Online]. Available: <https://api.semanticscholar.org/CorpusID:9465504>
- [23] J. Zhang, X. Tao, and H. Wang, "Outlier detection from large distributed databases," *World Wide Web*, vol. 17, pp. 539 – 568, 2013. [Online]. Available: <https://api.semanticscholar.org/CorpusID:13783878>
- [24] H. Wang and Y. Zhang, "Untraceable off-line electronic cash flow in e-commerce," *Proceedings 24th Australian Computer Science Conference. ACSC 2001*, pp. 191–198, 2001. [Online]. Available: <https://api.semanticscholar.org/CorpusID:1726261>
- [25] M. E. Kabir, H. Wang, and E. Bertino, "A conditional purpose-based access control model with dynamic roles," *Expert Syst. Appl.*, vol. 38, pp. 1482–1489, 2011. [Online]. Available: <https://api.semanticscholar.org/CorpusID:13760716>
- [26] H. Wang, X. Yi, E. Bertino, and L. Sun, "Protecting outsourced data in cloud computing through access management," *Concurrency and Computation: Practice and Experience*, vol. 28, pp. 600 – 615, 2016. [Online]. Available: <https://api.semanticscholar.org/CorpusID:39768338>
- [27] X. Sun, H. Wang, and J. Li, "Injecting purpose and trust into data anonymisation," *Proceedings of the 18th ACM conference on Information and knowledge management*, 2009. [Online]. Available: <https://api.semanticscholar.org/CorpusID:11922458>
- [28] M. Nasr, R. Shokri, and A. Houmansadr, "Comprehensive privacy analysis of deep learning: Passive and active white-box inference attacks against centralized and federated learning," in *2019 IEEE symposium on security and privacy (SP)*. IEEE, 2019, pp. 739–753.
- [29] L. Melis, C. Song, E. De Cristofaro, and V. Shmatikov, "Exploiting unintended feature leakage in collaborative learning," in *2019 IEEE symposium on security and privacy (SP)*. IEEE, 2019, pp. 691–706.
- [30] M. Abadi, A. Chu, I. Goodfellow, H. B. McMahan, I. Mironov, K. Talwar, and L. Zhang, "Deep learning with differential privacy," in *Proceedings of the 2016 ACM SIGSAC conference on computer and communications security*, 2016, pp. 308–318.
- [31] K. Bonawitz, V. Ivanov, B. Kreuter, A. Marcedone, H. B. McMahan, S. Patel, D. Ramage, A. Segal, and K. Seth, "Practical secure aggregation for privacy-preserving machine learning," in *proceedings of the 2017 ACM SIGSAC Conference on Computer and Communications Security*, 2017, pp. 1175–1191.
- [32] K. Wei, J. Li, M. Ding, C. Ma, H. H. Yang, F. Farokhi, S. Jin, T. Q. Quek, and H. V. Poor, "Federated learning with differential privacy: Algorithms and performance analysis," *IEEE transactions on information forensics and security*, vol. 15, pp. 3454–3469, 2020.
- [33] A. Triastcyn and B. Faltings, "Bayesian differential privacy for machine learning," in *International Conference on Machine Learning*. PMLR, 2020, pp. 9583–9592.
- [34] Y. Li, S. Wang, C.-Y. Chi, and T. Q. Quek, "Differentially private federated clustering over non-iid data," *IEEE Internet of Things Journal*, vol. 11, no. 4, pp. 6705–6721, 2023.
- [35] M. You, Y.-F. Ge, J. Yin, K. Wang, Z. Zheng, Y. Zhang, and H. Wang, "Akief: Adaptive knowledge inheritance evolutionary framework for dynamic privacy-preserving data publishing," *ACM Transactions on the Web*, 2025.
- [36] G. Andrew, O. Thakkar, B. McMahan, and S. Ramaswamy, "Differentially private learning with adaptive clipping," *Advances in Neural Information Processing Systems*, vol. 34, pp. 17 455–17 466, 2021.
- [37] L. Sun, J. Qian, and X. Chen, "Ldp-fl: Practical privacy aggregation in federated learning with local differential privacy," *arXiv preprint arXiv:2007.15789*, 2020.
- [38] J. Fu, Z. Chen, and X. Han, "Adap dp-fl: Differentially private federated learning with adaptive noise," in *2022 IEEE international conference on trust, security and privacy in computing and communications (Trust-Com)*. IEEE, 2022, pp. 656–663.
- [39] M. Alkhatami, Y.-F. Ge, and H. Wang, "A variation-based genetic algorithm for privacy-preserving data publishing," in *2024 11th International Conference on Machine Intelligence Theory and Applications (MiTA)*. IEEE, 2024, pp. 1–8.
- [40] F. Mireshghallah, H. A. Inan, M. Hasegawa, V. Rühle, T. Berg-Kirkpatrick, and R. Sim, "Privacy regularization: Joint privacy-utility optimization in language models," *arXiv preprint arXiv:2103.07567*, 2021.
- [41] Y.-F. Ge, H. Wang, J. Cao, Y. Zhang, and G. Kambourakis, "Federated genetic algorithm: Two-layer privacy-preserving trajectory data publishing," in *Proceedings of the genetic and evolutionary computation conference*, 2024, pp. 749–758.
- [42] Y. Shi, Y. Liu, K. Wei, L. Shen, X. Wang, and D. Tao, "Make landscape flatter in differentially private federated learning," in *Proceedings of the IEEE/CVF conference on computer vision and pattern recognition*, 2023, pp. 24 552–24 562.
- [43] L. Sun, H. Wang, J. Yong, and G. Wu, "Semantic access control for cloud computing based on e-healthcare," *Proceedings of the 2012 IEEE 16th International Conference on Computer Supported Cooperative Work in Design (CSCWD)*, pp. 512–518, 2012. [Online]. Available: <https://api.semanticscholar.org/CorpusID:7877201>
- [44] F. Zhang, Y. A. Wang, S. Liu, and H. Wang, "Decision-based evasion attacks on tree ensemble classifiers," *World Wide Web*, vol. 23, pp. 2957 – 2977, 2020. [Online]. Available: <https://api.semanticscholar.org/CorpusID:215819534>
- [45] H. Wang, Y. Zhang, and J. Cao, "Ubiquitous computing environments and its usage access control," in *Scalable Information Systems*, 2006. [Online]. Available: <https://api.semanticscholar.org/CorpusID:7988205>